

2021 年中终结

张鼎

武汉大学计算机学院

2021 年 07 月 29 日



① 回顾

② 展望



Outline

① 回顾

② 展望



回顾

- 冬奥气象

- ▷ 任务描述：输入一句话，返回需要提取的要素。（比去年的气象项目简单）
- ▷ 要素包括：天气、时间、地点。
- ▷ 经过测试，天气、地点要素由模型提取，时间由正则表达式提取效果最好。

回顾

- 冬奥气象

- ▷ 使用 Docker 打包为镜像，一键部署、维护。
- ▷ 使用 buildx 打包镜像，可以使用 `--mount=type=bind,target=` 来避免临时文件被打包到最终镜像中，缩小镜像大小。

```
#-----  
  
RUN --mount=type=bind,target=/build_tmp/ \  
... pip3 install --no-cache-dir -i https://mirrors.aliyun.com/pypi/simple/ /build_tmp/  
... pip_package_file/tensorflow-1.15.0-cp37-cp37m-manylinux2010_x86_64.whl && \  
... pip3 install --no-cache-dir -i https://mirrors.aliyun.com/pypi/simple/ /build_tmp/  
... pip_package_file/torch-1.9.0+cu111-cp37-cp37m-linux_x86_64.whl && \  
... pip3 install --no-cache-dir -i https://mirrors.aliyun.com/pypi/simple/ /build_tmp/  
... pip_package_file/torchvision-0.10.0+cu111-cp37-cp37m-linux_x86_64.whl  
  
RUN --mount=type=bind,target=/build_tmp/ \  
... pip3 install --no-cache-dir -i https://mirrors.aliyun.com/pypi/simple/ -r /build_tmp/  
... requirements.txt && \  
... pip3 install --no-cache-dir -i https://mirrors.aliyun.com/pypi/simple/ uwsgi
```

回顾

● 冬奥气象

- ▷ 使用 Docker 打包为镜像，一键部署、维护。
- ▷ 使用 buildx 打包镜像，可以使用 `--mount=type=bind,target=` 来避免临时文件被打包到最终镜像中，缩小镜像大小。
- ▷ 可以在 `docker-compose.yml` 文件中写入 `deploy/devices` 配置来使得容器可以使用显卡。

```
deploy:
  resources:
    reservations:
      devices:
        - driver: nvidia
          capabilities: ["compute", "utility"]
          device_ids: ["0"]
```

回顾

● 冬奥气象

- ▷ 使用 Docker 打包为镜像，一键部署、维护。
- ▷ 使用 buildx 打包镜像，可以使用 `--mount=type=bind,target=` 来避免临时文件被打包到最终镜像中，缩小镜像大小。
- ▷ 可以在 `docker-compose.yml` 文件中写入 `deploy/devices` 配置来使得容器可以使用显卡。
- ▷ 使用 `cuda` 镜像请优先使用 `nvcr.io/nvidia/cuda` 地址（更新、更快）。

回顾

- 冬奥气象
- 服务器监控系统
 - ▷ 架构设计/选好了，但是部署的时候一直出问题，最后并没有实际部署监控系统。
 - ▷ 通常的方法（直接使用 Prometheus）在我们学院的网络环境下完全不可用。子网之间乱隔离。（每日吐槽学院网络 ✓）

回顾 - 服务器监控 - 架构

★ 监控、报警

- ◇ Prometheus
- ◇ Alertmanager

★ 存储、查询、规则：

- ◇ Thanos Receive
- ◇ Thanos Ruler
- ◇ Thanos Store Gateway
- ◇ Thanos Compact
- ◇ Thanos Query
- ◇ Thanos Query Frontend

★ 可视化：

- ◇ Grafana

回顾

- 冬奥气象
- 服务器监控系统
- 服务器升级（内存和显卡）
 - ▷ 内存安装的插槽位置非常重要！被苏州超集的工程师给坑了。最开始工程师装的是错的，导致服务器运行非常不稳定。
 - ▷ 内存安装的位置一定要按照主板用户手册上的推荐配置表来。

回顾 - 服务器升级 - 内存安装表

Chapter 3: Maintenance and Component Installation

Memory Population Table for the X11DP Motherboard w/24 DIMM Slots Onboard	
CPU/DIMMs	Memory Population Sequence
1 CPU & 1 DIMM	CPU1: P1-DIMMA1
1 CPU & 2 DIMMs	CPU1: P1-DIMMA1P1-DIMMD1
1 CPU & 3 DIMMs	CPU1: P1-DIMMC1P1-DIMMB1P1-DIMMA1
1 CPU & 4 DIMMs	CPU1: P1-DIMMB1P1-DIMMA1P1-DIMMD1P1-DIMME1
1 CPU & 5 DIMMs*	CPU1: P1-DIMMC1P1-DIMMB1P1-DIMMA1P1-DIMMD1P1-DIMME1
1 CPU & 6 DIMM	CPU1: P1-DIMMC1P1-DIMMB1P1-DIMMA1P1-DIMMD1P1-DIMME1P1-DIMMF1
1 CPU & 7 DIMMs*	CPU1: P1-DIMMB1P1-DIMMB2P1-DIMMA1P1-DIMMA2P1-DIMMD1P1-DIMME1P1-DIMMF1
1 CPU & 8 DIMMs	CPU1: P1-DIMMB1P1-DIMMB2P1-DIMMA1P1-DIMMA2P1-DIMMD2P1-DIMMD1P1-DIMME1P1-DIMMF1
1 CPU & 9 DIMMs*	CPU1: P1-DIMMC1P1-DIMMC2P1-DIMMB1P1-DIMMB2P1-DIMMA1P1-DIMMA2P1-DIMMD1P1-DIMME1P1-DIMMF1
1 CPU & 10 DIMMs*	CPU1: P1-DIMMC1P1-DIMMB1P1-DIMMB2P1-DIMMA1P1-DIMMA2P1-DIMMD2P1-DIMMD1P1-DIMME1P1-DIMMF1
1 CPU & 11 DIMMs*	CPU1: P1-DIMMC1P1-DIMMC2P1-DIMMB1P1-DIMMB2P1-DIMMA1P1-DIMMA2P1-DIMMD2P1-DIMMD1P1-DIMME1P1-DIMMF1
1 CPU & 12 DIMMs	CPU1: P1-DIMMC1P1-DIMMC2P1-DIMMB1P1-DIMMB2P1-DIMMA1P1-DIMMA2P1-DIMMD2P1-DIMMD1P1-DIMME1P1-DIMMF1
2 CPUs & 2 DIMMs	CPU1: P1-DIMMA1 CPU2: P2-DIMMA1
2 CPUs & 4 DIMMs	CPU1: P1-DIMMA1P1-DIMMD1 CPU2: P2-DIMMA1P2-DIMMD1
2 CPUs & 6 DIMMs	CPU1: P1-DIMMC1P1-DIMMB1P1-DIMMA1 CPU2: P2-DIMMC1P2-DIMMB1P2-DIMMA1
2 CPUs & 8 DIMMs	CPU1: P1-DIMMB1P1-DIMMA1P1-DIMMD1P1-DIMME1 CPU2: P2-DIMMB1P2-DIMMA1P2-DIMMD1P2-DIMME1
2 CPUs & 10 DIMMs	CPU1: P1-DIMMC1P1-DIMMB1P1-DIMMA1P1-DIMMD1P1-DIMME1P1-DIMMF1 CPU2: P2-DIMMC1P2-DIMMB1P2-DIMMA1P2-DIMMD1P2-DIMME1P2-DIMMF1
2 CPUs & 12 DIMMs	CPU1: P1-DIMMC1P1-DIMMB1P1-DIMMA1P1-DIMMD1P1-DIMME1P1-DIMMF1 CPU2: P2-DIMMC1P2-DIMMB1P2-DIMMA1P2-DIMMD1P2-DIMME1P2-DIMMF1
2 CPUs & 14 DIMMs	CPU1: P1-DIMMC1P1-DIMMB1P1-DIMMA1P1-DIMMD1P1-DIMME1P1-DIMMF1 CPU2: P2-DIMMC1P2-DIMMB1P2-DIMMA1P2-DIMMD1P2-DIMME1P2-DIMMF1
2 CPUs & 16 DIMMs	CPU1: P1-DIMMB1P1-DIMMB2P1-DIMMA1P1-DIMMA2P1-DIMMD2P1-DIMMD1P1-DIMME1P1-DIMMF1 CPU2: P2-DIMMB1P2-DIMMB2P2-DIMMA1P2-DIMMA2P2-DIMMD2P2-DIMMD1P2-DIMME1P2-DIMMF1
2 CPUs & 18 DIMMs	CPU1: P1-DIMMC1P1-DIMMC2P1-DIMMB1P1-DIMMB2P1-DIMMA1P1-DIMMA2P1-DIMMD2P1-DIMMD1P1-DIMME1P1-DIMMF1 CPU2: P2-DIMMC1P2-DIMMC2P2-DIMMB1P2-DIMMB2P2-DIMMA1P2-DIMMA2P2-DIMMD2P2-DIMMD1P2-DIMME1P2-DIMMF1
2 CPUs & 20 DIMMs	CPU1: P1-DIMMC1P1-DIMMC2P1-DIMMB1P1-DIMMB2P1-DIMMA1P1-DIMMA2P1-DIMMD2P1-DIMMD1P1-DIMME1P1-DIMMF1 CPU2: P2-DIMMC1P2-DIMMC2P2-DIMMB1P2-DIMMB2P2-DIMMA1P2-DIMMA2P2-DIMMD2P2-DIMMD1P2-DIMME1P2-DIMMF1
2 CPUs & 22 DIMMs*	CPU1: P1-DIMMC1P1-DIMMC2P1-DIMMB1P1-DIMMB2P1-DIMMA1P1-DIMMA2P1-DIMMD2P1-DIMMD1P1-DIMME1P1-DIMMF1 CPU2: P2-DIMMC1P2-DIMMC2P2-DIMMB1P2-DIMMB2P2-DIMMA1P2-DIMMA2P2-DIMMD2P2-DIMMD1P2-DIMME1P2-DIMMF1
2 CPUs & 24 DIMMs	CPU1: P1-DIMMC1P1-DIMMC2P1-DIMMB1P1-DIMMB2P1-DIMMA1P1-DIMMA2P1-DIMMD2P1-DIMMD1P1-DIMME1P1-DIMMF1 CPU2: P2-DIMMC1P2-DIMMC2P2-DIMMB1P2-DIMMB2P2-DIMMA1P2-DIMMA2P2-DIMMD2P2-DIMMD1P2-DIMME1P2-DIMMF1

*Unbalanced, not recommended.

回顾

- 冬奥气象
- 服务器监控系统
- 服务器升级（内存和显卡）
- 服务器显卡微调
 - ▷ 升级显卡 vbios

回顾

- 冬奥气象
- 服务器监控系统
- 服务器升级（内存和显卡）
- 服务器显卡微调
 - ▷ 升级显卡 vbios
 - ▷ 打开显卡的 persistence mode，防止显卡深度休眠。显著加快显卡 ready 时间（显卡越多越明显），对功耗基本没影响。

回顾

- 冬奥气象
- 服务器监控系统
- 服务器升级（内存和显卡）
- 服务器显卡微调
- 刷题 & 背八股文
- 财务交接
- 项目预算使用情况统计
- 调账 & 调预算

Outline

① 回顾

② 展望



展望

- 找工作
 - ▷ 刷题 & 背八股文
- 完成毕业论文
 - ▷ 以气象项目为基础。目前还没有什么好的想法/还没有认真的想...
 - ▷ (结合知识图谱?)



Thanks for listening!